

Experiment 3: Thresholding and contours

Adrian F. Clark

This experiment explores the thresholding capabilities of OpenCV and introduces its contour-processing routines through the example of counting the number of dots on several dice.

Contents

1	Introduction	2
2	The program	2
3	Thresholding	4
4	Processing contours	5
5	Counting the dots	6
6	Concluding remarks	7

1 Introduction

This experiment is based on a tutorial on contours using OpenCV which first touches on its thresholding and then its contour functionalities; alas that tutorial has now disappeared from the Web. Thresholding is discussed in detail Chapter 5 of the lecture notes but then we considered region labelling rather than contours. Contours are ultimately more powerful but the code to implement it is significantly more complicated and the way in which it is used from OpenCV is definitely more difficult.

The particular image we shall use for this experiment is shown in Figure 1. The aim of this experiment is to isolate each of the visible die faces by thresholding and then count the number of dots on each die's face.



Figure 1: Example image for this experiment

2 The program

This experiment uses the following program:

```
#!/usr/bin/env python3
"contours -- demo of OpenCV's contour-processing capabilities"
import sys, cv2, numpy

# Set-up.
thresholding = "vanilla"
thresholding = "adaptive"
thresholding = "otsu"

contouring = "vanilla"
contouring = "tree"
```

```

# Handle the command line.
if len (sys.argv) < 3:
    print ("Usage:", sys.argv[0], "<image> <threshold>", file=sys.stderr)
    sys.exit (1)
im = cv2.imread (sys.argv[1])
t = int (sys.argv[2])

# Convert to greyscale and reduce noise a little.
grey = cv2.cvtColor (im, cv2.COLOR_BGR2GRAY)
blur = cv2.GaussianBlur (grey, (5, 5), 0)

# Threshold the image.
if thresholding == "adaptive":
    binary = cv2.adaptiveThreshold (blur, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
                                   cv2.THRESH_BINARY, 11, 2)
elif thresholding == "otsu":
    t, binary = cv2.threshold (blur, 0, 255, cv2.THRESH_BINARY
                              + cv2.THRESH_OTSU)
    print ("Otsu threshold is", t)
else:
    t, binary = cv2.threshold (blur, t, 255, cv2.THRESH_BINARY)

# Remove small features in the binary image using a morphological close.
kernel = numpy.ones ((9,9), numpy.uint8)
binary = cv2.erode (binary, kernel, iterations=1)

# Find contours.
if contouring == "tree":
    # Find internal and external contours.
    contours, hierarchy = cv2.findContours (binary, cv2.RETR_TREE,
                                           cv2.CHAIN_APPROX_SIMPLE)

    # Count the number of dots on the dice faces. We do this by iterating
    # through hierarchy[0], first to find the indices of dice contours,
    # then a second time to find dot contours.
    dice = [] # list of dice contours
    dots = [] # list of dot contours

    # Find dice contours, drawing the contours as we process them.
    for (i, c) in enumerate(hierarchy[0]):
        if c[3] == -1:
            dice.append (i)
            cv2.drawContours (im, contours, i, (0, 0, 255), 5)

    # Find dot contours, drawing them as we process them.
    for (i, c) in enumerate(hierarchy[0]):
        if c[3] in dice:
            dots.append (i)

```

```

        cv2.drawContours (im, contours, i, (0, 255, 0), 5)

    # Report the total number of dots found.
    print ("Total die roll:", len (dots))

else:
    # Find external contours only.
    contours, junk = cv2.findContours (binary, cv2.RETR_EXTERNAL,
                                       cv2.CHAIN_APPROX_SIMPLE)

    # Draw contours over original image.
    cv2.drawContours (im, contours, -1, (0, 0, 255), 5)

    # Print a table of the contours and their sizes.
    print ("Found %d objects." % len(contours))
    for (i, c) in enumerate(contours):
        print ("\tSize of contour %d: %d" % (i, len(c)))

# Display the result.
cv2.namedWindow (sys.argv[0], cv2.WINDOW_NORMAL)
ny, nx, nc = im.shape
cv2.resizeWindow (sys.argv[0], nx//2, ny//2)
cv2.imshow (sys.argv[0], im)
cv2.waitKey (0)

```

Rather than type it in, you should download [a zip-file](#) containing it and the example image of Figure 1. Like most programs that work purely by calling OpenCV routines, it is fairly short. It works by reading in an image and blurring it slightly with a Gaussian-shaped mask. It then thresholds the image using an OpenCV function. The resulting binary image is passed to OpenCV's contour-finding function, which essentially returns a list of the contours found. These are printed out and then drawn onto the image for display. However, there are alternative ways of doing the thresholding and contour-processing, controlled by the variables `thresholding` and `contouring` respectively.

You should run the program unchanged on the supplied image and see how well it works. You invoke it with two arguments, the name of the file containing the image to be processed and the value of the threshold; try using 200 first and then see what effect changing it has.

3 Thresholding

You will see that the routine `cv2.threshold` accepts several arguments. Look at its documentation on the Web so that you understand what they all do. While you are doing that, also make sure you understand what the function returns: the thresholded image is only one of its outputs.

One characteristic of the dice image is that its background is not uniform but appears lighter towards the upper right corner. That is not a problem in this case because the upper surfaces of the dice appear so white; but in many tasks, this variation in the background may become problematic. For this reason, OpenCV provides an *adaptive* thresholder, where the threshold relates to regions of the image rather than being global. This routine is used in the following way:

```
binary = cv2.adaptiveThreshold (blur, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
                                cv2.THRESH_BINARY, 11, 2)
```

where the last two arguments can be tuned to make the routine perform well — you should read up on them too, to help you work out good values. Change the value of thresholding to "adaptive" and see if the adaptive threshold works as well. Can you identify any difference in the run-times of the two versions? (You can use Unix's `time` command to do this.)

One irritation of both the standard and adaptive thresholders is that the value of the threshold has to be provided by you on the command line. The underlying aim of computer vision is for the entire process to be automated, so it makes sense to work out the value of the threshold from the image if at all possible. Lectures described the use of Otsu's approach, which works well when the image is bimodal (*i.e.*, its histogram has two peaks).

You can use Otsu's method with `cv2.threshold` too, using a call like

```
thresh, binary = cv2.threshold (blur, 0, 255,
                                cv2.THRESH_BINARY + cv2.THRESH_OTSU)
```

This has an additional option in the last argument and uses zero for the threshold argument; together, these tell `cv2.threshold` to compute the threshold value using Otsu's method. The first argument returned from the call, stored in `thresh` here, is the threshold value computed.

Change the value of thresholding to "otsu" and try it — and remember that you have to provide a value of the threshold on the command line even if it isn't used. Does Otsu's method work well in this case? What rules of thumb were mentioned in the lecture regarding its use? Are they met here?

4 Processing contours

Let us now turn our attention to finding and processing contours in the image. We shall start with contouring set to "vanilla". The call to `cv2.findContours` passes in three arguments and receives two outputs. The first parameter in the call is easy enough, the image in which the contours are to be found. This image should be binary, with the objects for which contours are to be found in white and a black background.

The second argument is a constant indicating what kind of contours are to be found. As we are currently interested in the main

objects in the image, we want only contours around the outermost edges of objects and so we pass in `cv2.RETR_EXTERNAL`. If we required more information, such as the contours of the dots marked on a die's face, then we would use another parameter such as `cv2.RETR_TREE` or `cv2.RETR_CCMP`. (We shall come back to this shortly.) The last parameter tells `cv2.findContours` whether or not it should simplify the contours; we use `cv2.CHAIN_APPROX_SIMPLE`, which tells it to simplify by using line segments when it can — that saves memory and computation time here.

The first return value is a list of numpy arrays, each holding the points for a single contour in the image. The other return value is a numpy array that contains hierarchy information about the contours, again not of interest to us here.

Knowing this, the remainder of the program should be straightforward to understand.

5 Counting the dots

The program initially finds only the external contours of objects. Changing the value of `contouring` to "tree" makes the program also find internal contours. You will see that it does this by changing the third argument from `cv2.RETR_EXTERNAL` to `cv2.RETR_TREE` — and we can use this to count the number of dots on each die's face.

When using `cv2.RETR_TREE`, the contours are arranged in a hierarchy, with the outermost contours for each object at the top. Moving down the hierarchy, each new level of contours represents the next innermost contour for each object. We obtain the contour hierarchies via the second return value from the `cv2.findContours` call:

```
contours, hierarchy = cv2.findContours (binary, cv2.RETR_TREE,
                                       cv2.CHAIN_APPROX_SIMPLE)
```

`hierarchy` is a three-dimensional numpy array, with one row, 36 columns, and a "depth" of 4 from the test image in Figure 1. The 36 columns correspond to the contours found by the method — there are 36 contours rather than seven because `cv2.RETR_TREE` tells the routine to return internal contours as well as external ones. Column zero corresponds to the first contour, column one the second, and so on.

Each of the columns has a four-element array of integers, representing indices of other contours, according to the scheme

```
[next, previous, firstChild, parent]
```

The next index refers to the next contour in this contour's hierarchy level, while the previous index refers to the previous contour in this contour's hierarchy level. The `firstChild` index refers to the first contour that is contained inside this contour. The `parent` index refers to the contour containing this contour. In all cases, a value of `-1` indicates that there is no linked contour.

Things become clearer when different contours in the hierarchy are displayed in different colours. The program displays external

contours, those at the top of the hierarchy, in red. The next level of contours, which should be the dots on the faces, are outlined in green. You might try to draw the innermost contours, representing some lost paint in one of the dots in the central die, in blue.

How does the program use this information to count the number of dots? Does the program yield the correct count of dots?

6 Concluding remarks

Before putting this experiment to one side and doing something more interesting, you might like to spend a few minutes thinking about how well the software explored in this experiment will work on images other than the one you have used for testing. To do that, you might reflect on the properties of this image:

- it is monochrome rather than colour;
- the lighting is fairly flat, though not uniform across the field of view;
- there is very little noise in the image;
- the dots on the dice's faces are very dark;
- imperfections in the dots affect the contours that are found.

In fact, you might look for images that break each of these properties in turn and see how well the software handles it.